

Waliji Experiments

Mark Stoehr

September 4, 2012

Contents

1	Notation	1
2	Reproduce Waliji Parts	2
2.1	Sample work through	2
2.2	Clustering over many parts	11
2.2.1	Cluster Mapping	20

1 Notation

- π_k is a mixture weight
- e is an index for the edge orientation
- t, s are time indices
- f is a frequency band
- $X \in \{0, 1\}^{|\{e\}|, |\{f\}|, |\{s\}|}$ is the output of our edge detector ranging over edge locations $|\{e\}|$, frequency bands $|\{f\}|$, and time indices $|\{s\}|$
- $X_{e,f,s}$ is a given binary variable for whether an edge is detected at (e, f, s) in the collection denoted by X and $X_{i,e,f,s}$ is functions in the same role for a data point X_i .
- X denotes the random variable, x denotes the realized value
- $[n] = \{1, \dots, n\}$

2 Reproduce Waliji Parts

Reproducing the work from the Masters thesis

2.1 Sample work through

```
import matplotlib.pyplot as plt
import numpy as np
s = np.load('../data/1s.npy')
t = np.arange(len(s))/16.
plt.plot(t,s)
plt.xlabel('time (ms)')
plt.savefig('1s.png')
```




Figure 1: Example utterance waveform

We then turn this into a spectrogram. There are two ways of doing this the first is to use the same signal processing the Alexey Koloydenko and

Partha Niyogi developed

```
import matplotlib.pyplot as plt
import numpy as np
S = np.load('../data/1S.npy')
plt.imshow(S[::-1],interpolation='nearest',aspect=3)
plt.xticks(np.arange(7)*100,tuple(str(i*100/5) for i in xrange(7)))
# since 8*16000 (sample rate)/512 (fft length) = 250
plt.yticks(np.arange(6)*8,tuple(str(i*250) for i in xrange(6))[:,::-1])
plt.xlabel('time (ms)')
plt.ylabel('freq (Hz)')
plt.savefig('1S.png')
```




Figure 2: Example utterance Spectrogram

We can also redo this entirely as the mel spectrogram. More on that later.

We aim to extract edge features. The edge features we compute are done with binary masks of the following sort

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
edge_orientations = np.load('../data/edge_orientations.npy')
fig = plt.figure()
for i in xrange(8):
    plt.subplot(4,2,i+1)
    mask_mat = np.zeros((2,2))
    y_coord, x_coord = edge_orientations[i].astype(int)
    y = (y_coord + 1)/2
    x = (x_coord + 1)/2
    if y_coord == 0:
        plus_one_locs = [[0,1],[x,x]]
        minus_one_locs = [[0,1],[(x + 1) % 2, (x+1) % 2]]
    elif x_coord == 0:
        plus_one_locs = [[y,y],[0,1]]
        minus_one_locs = [[(y+1) % 2, (y+1) % 2], [0,1]]
    else:
        plus_one_locs = [[y],[x]]
        minus_one_locs = [[(y+1) % 2], [(x+1) % 2]]
    mask_mat[plus_one_locs] = 1
    mask_mat[minus_one_locs] = -1
    plt.imshow(mask_mat, interpolation='nearest',
               cmap=cm.bone)
    plt.title('Edge orientation %d' % (i+1))
    frame = plt.gca()
    frame.axes.get_xaxis().set_visible(False)
    frame.axes.get_yaxis().set_visible(False)

plt.subplots_adjust(hspace=.5)
plt.savefig('edge_orientations_pic.png')
```

These are represented in the following picture 3.

Applying those filters to the image and then thresholding the outputs we get the following edge map representation:

```
import matplotlib.pyplot as plt
```

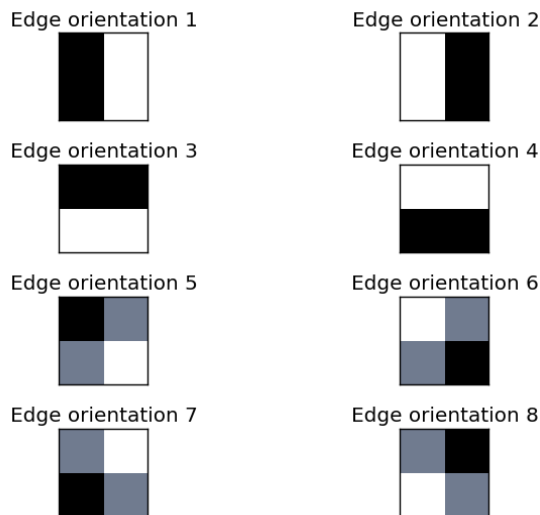


Figure 3: Edge filters

```
import matplotlib.cm as cm
import numpy as np
E = np.load('../data/1E.npy')
plt.imshow(E[::-1], interpolation='nearest',
           cmap=cm.bone)
plt.title('Edge Map Representation')
plt.xticks(np.arange(7)*100, tuple(str(i*100/5) for i in xrange(7)))
# since 8*16000 (sample rate)/512 (fft length) = 250
plt.yticks(np.arange(8)*48 + 24, tuple(str(i+1) for i in xrange(8))[::-1])
plt.xlabel('time (ms)')
plt.ylabel('freq (Hz)')
plt.savefig('edge_map.png')
```

In 4 the light areas indicate where edges have been detected. Also bear in mind that we spread the edges. Since there are eight orientations to the edges each edge filter produces a separate copy of the spectrogram. We stacked these representations onto the same figure and the y-axis consists of blocks of these representation. Increasing along the y-axis are the frequency locations.

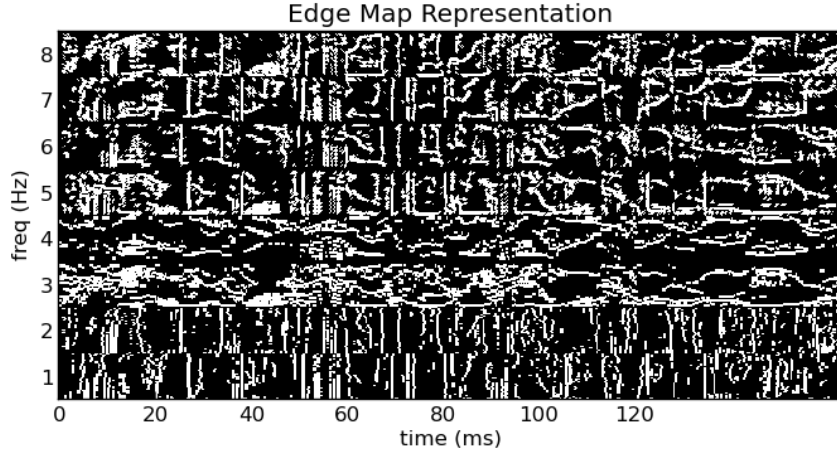


Figure 4: Edge Map Representation

We then extract features related to 6 by 6 square patches in the spectrogram. These features are associated with an 8 by 5 by 5 patch of the edge map, where we have a 5 by 5 patch from each edge map associated with a particular edge orientation. We only consider such patches if the number of edges is above a certain threshold: chosen to be the 90th percentile following Waliji.

In the particular example above we have 21120 total patches and the distribution over the number of edges is given in 5 and the cutoff point is at about 70 edges. Each patch can have potentially 200 edges in it (although its impossible for a signal to have that many edges).

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
bp_all = np.load('../data/1bp_all.npy')
plt.close()
```

```
plt.hist(bp_all.sum(1).sum(1))
plt.title('Histogram over number of edges in given patches')
plt.savefig('edge_count_patches_histogram.png')
plt.close()
```




Figure 5: Patch Edge Count Histogram

We see that the number of edges in a patch is approximately normal. We can then look at the spectrogram patches that are associated with these edgemap patches.

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
from sklearn import mixture

spec_patch = np.load('../data/1spec_patch.npy')
spec_patch_flat = spec_patch.reshape(spec_patch.shape[0], 6*6)
```

```

for i in [1,2,3,5,8,13,21]:
    clf = mixture.GMM(n_components=i,n_init=10)
    clf.fit(spec_patch_flat)
    num_rows = i/3+1
    if i < 3:
        num_cols = i+1
    else:
        num_cols = 3
    fig = plt.figure()
    plt.title('Mixture Components')
    for j in xrange(i):
        plt.subplot(num_rows,num_cols,j+1)
        plt.imshow(clf.means_[j].reshape(6,6),
                    cmap=cm.bone)
        frame = plt.gca()
        frame.axes.get_xaxis().set_visible(False)
        frame.axes.get_yaxis().set_visible(False)
    plt.savefig('spec_patch_clustersGMM%d.png' % i)

import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
from sklearn.cluster import MeanShift, estimate_bandwidth

spec_patch = np.load('../data/1spec_patch.npy')
X = spec_patch.reshape(spec_patch.shape[0],6*6)
for bandwidth in [16,8,4,2,1,.5,.25,.125]:
    ms = MeanShift(bandwidth=bandwidth, bin_seeding=True)
    ms.fit(X)
    labels = ms.labels_
    cluster_centers = ms.cluster_centers_
    labels_unique = np.unique(labels)
    n_clusters_ = len(labels_unique)
    print "number of estimated clusters : %d" % n_clusters_
    num_rows = n_clusters_/3+1
    if n_clusters_ < 3:
        num_cols = n_clusters_+1
    else:
        num_cols = 3

```





Figure 6: Spectrogram Patch Mean

```
fig = plt.figure()
plt.title('Mixture Components')
for j in xrange(n_clusters_):
    plt.subplot(num_rows,num_cols,j+1)
    plt.imshow(cluster_centers[j].reshape(6,6),
               cmap=cm.bone)
    frame = plt.gca()
    frame.axes.get_xaxis().set_visible(False)
    frame.axes.get_yaxis().set_visible(False)
plt.savefig('spec_patch_clustersMeanShift%d%2f.png' % (n_clusters_,bandwidth))
```

Now, we see that these pick up on the edge structure of the spectrogram quite nice, just as we would expect. Our next question is where these are coming from in the spectrogram, we show this in ??, the blue denotes areas where no patches have been extracted, the colors give a sense of how the extraction process picks up on loud and quiet parts of the spectrogram.




Figure 7: Spectrogram Patch GMM - 2 Clusters

```
import matplotlib.pyplot as plt
import numpy as np
S = np.load('../data/1S.npy')
spec_patch_ones = np.load('../data/1spec_patch_ones.npy')
S -= S.min()
S /= S.max()
S *= spec_patch_ones
plt.imshow(S[::1], interpolation='nearest', aspect=3)
plt.xticks(np.arange(7)*100, tuple(str(i*100/5) for i in xrange(7)))
# since 8*16000 (sample rate)/512 (fft length) = 250
plt.yticks(np.arange(6)*8, tuple(str(i*250) for i in xrange(6))[::-1])
plt.xlabel('time (ms)')
plt.ylabel('freq (Hz)')
plt.savefig('1S_spec_patch.png')
```



Figure 8: Spectrogram Patch GMM - 3 Clusters

We then consider what happens when we apply the bernoulli clustering algorithm (with EM) to the binary features. A single utterance gives us a fair sampling of the patches.

2.2 Clustering over many parts

From 28 utterances we extract a total of 50,000 patches (again these are in the top 90th percentile of edges for the utterance) and we do clustering over these. The clustering is over the binary edge maps for the patches. We clustered with 10, 20, 30, 50, 80, and 100 cluster centers to get a sense of what different numbers of clusters mean.

Bernoulli Mixture models in high dimensions estimate a probability mass function of the form




Figure 9: Spectrogram Patch GMM - 5 Clusters

$$\mathbb{P}(X) = \sum_{k=1}^K \pi_k \prod_{e,f,s} X_{e,f,s}^{p_{k,e,f,s}} (1 - X_{e,f,s})^{p_{k,e,f,s}}$$

where K is the number of components in the mixture so k is the component identity, the index e ranging over $1, 2, \dots, 8$ represents the edge orientation, f is the frequency band, and s is the time. In the case of patches f is not absolute, but instead is relative to the lowest frequency band in the patch, as we do not treat patches extracted from the lower part of the spectrogram as being different from those extracted from a higher part of the spectrogram. s is also not an absolute time either but is also relative to when the patch begins. The patches range over five edge frequency bands and five edge time blocks. So $f \in [5]$ and $s \in [5]$.

We perform clustering using the EM algorithm. When running the EM algorithm we compute ‘cluster affinities’ $A_{i,k}$ - which are for each mixture component k of our mixture model we compute the probability that a given




Figure 10: Spectrogram Patch GMM - 8 Clusters

data point, X_i , was generated by that mixture component. More formally, a given mixture model is specified by associated a latent ‘label’ Z_i with each datapoint X_i . This label $Z_i \in [K]$ where K is the number of components. We model the binary variabls that make up a data point X_i as conditionally independent bernoulli trials given the label Z_i . The affinity is $\mathbb{P}(Z_i = k \mid X_i)$

Our formula is

$$A_{i,k} = \frac{\pi_k \prod_{e,f,s} X_{i,e,f,s}^{p_{k,e,f,s}} (1 - X_{i,e,f,s})^{1-p_{k,e,f,s}}}{\sum_{k'=1}^K \pi_{k'} \prod_{e,f,s} X_{i,e,f,s}^{p_{k',e,f,s}} (1 - X_{i,e,f,s})^{1-p_{k',e,f,s}}}$$

We find that in our bernoulli mixture model that these affinities tend to be highly degenerate and either very close to 0 or very close to 1.

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from matplotlib import rc
```




Figure 11: Spectrogram Patch GMM - 13 Clusters

```
import numpy as np
rc('text', usetex=True)
fig = plt.figure()
plt.title('CDFs over  $\max_k P(Z_i=k \mid X_i = x_i)$ ')
for mix_id, num_mix in enumerate([10,20,30,50,80,100]):
    plt.subplot(2,3,mix_id+1)
    sorted_affinities= np.sort(np.load('../data/bm_affinities%d.npy' % num_mix).astype(float))
    plt.plot(sorted_affinities,
             np.arange(sorted_affinities.shape[0])/float(len(sorted_affinities))
            )
    plt.title('Empirical cdf \nfor %d components' % num_mix)

plt.subplots_adjust(wspace=.5,hspace=.5)

plt.savefig('mixture_affinities_cdfs.png')
```




Figure 12: Spectrogram Patch GMM - 21 Clusters

```
plt.close()
```

We then see in 14 that essentially every data point is assigned to single component with overwhelming probability. The curves in those plots are of the function

$$\left(\tau, \frac{|\{i \mid \max_k A_{i,k} < \tau\}|}{n} \right)$$

which is the empirical cumulative distribution function of the maximum estimated affinities for the data points.

The consequence of the fact that the affinities are degenerate is that essentially the bernoulli mixture modeling provides a hard clustering of the extracted patches.

We can then examine the cluster centers given by the templates, and, more informatively, the affinities induce a clustering on the spectrogram patches that gave rise to the bernoulli edge map features.




Figure 13: Example utterance Spectrogram - Edge Patch Locations

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from matplotlib import rc
import numpy as np
rc('text', usetex=True)
for num_mix in [10,20,30,50,80,100]:
    spec_avg_parts = np.load('../data/spec_avg_parts%d.npy' % num_mix)
    fig = plt.figure()
    for i in xrange(num_mix):
        plt.subplot((num_mix-1)/4+1,min(num_mix+1,4),i+1)
        plt.imshow(spec_avg_parts[i],cmap=cm.bone,interpolation='nearest')
        frame = plt.gca()
        frame.axes.get_xaxis().set_visible(False)
        frame.axes.get_yaxis().set_visible(False)
    plt.savefig('spec_avg_parts_%d.png' % num_mix)
```

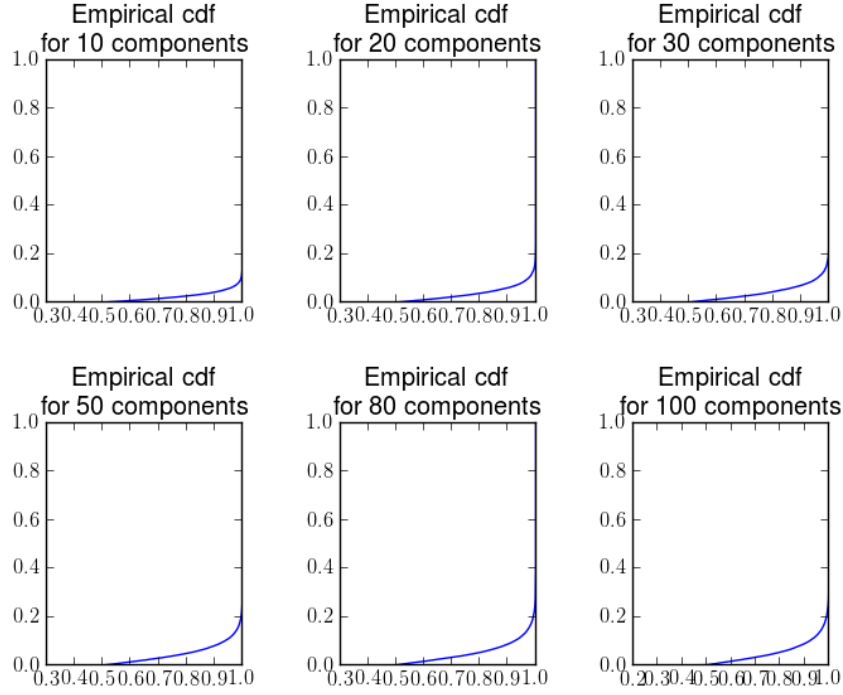



Figure 14: Empirical CDFs for the Max Component Affinities

We also want to see what the parts are when we use a single utterance. There is a noticeable lack of parts for vertical objects, whereas this was observed in Waliji’s experiment based on the figures provided.

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
from matplotlib import rc
import numpy as np
num_mix = 20
spec_avg_parts = np.load('../data/1spec_avg_parts%d.npy'
                          % num_mix)

fig = plt.figure()
for i in xrange(num_mix):
    plt.subplot((num_mix-1)/4+1,min(num_mix+1,4),i+1)
    plt.imshow(spec_avg_parts[i],cmap=cm.bone,interpolation='nearest')
    frame = plt.gca()
```




Figure 15: Spectrogram Part Clusters - 10 components

```
frame.axes.get_xaxis().set_visible(False)
frame.axes.get_yaxis().set_visible(False)
plt.savefig('lspec_avg_parts_%d.png' % num_mix)
```

Even with only 2600 patches we see roughly the same patterns as we do with many more patches as in ??.

The apparent lack of time-edges, that is, edges where the gradient is zero across frequency bands but large between successive time points, is likely going to be problematic for classification. A useful statistic to compute is in what clusters are those vertical edges being assigned? Additionally, one question is whether vertical edge structure preserved under the feature map? There are a couple of approaches to answering these questions:

- cluster assignment inspection: among those patches that we declare non-background, how does the number of time edges correlate with cluster assignment. Are time-edge heavy patches assigned to many




Figure 16: Spectrogram Part Clusters - 20 components

different clusters or are they assigned to a particular set.

- Time edges usually indicate broad-band noise (this is also something to verify) about the speech signal, we should note when those occur versus when they don't how that correlates to the patches
- we should see if the patches with the most time-edges are unfairly penalized in that they tend to have fewer edges of the other types, and hence get ignored
- we should also see what the classification performance of using these feature maps are, in particular, does the presence of time edges correlate with our misclassifications?




Figure 17: Spectrogram Part Clusters - 30 components

2.2.1 Cluster Mapping

In order to do the cluster mapping process and get a proper feature map, we take as input an edge map $E(t, f, e) \in \{0, 1\}$ and we construct a feature map $\Phi(t, f) \in \{-1, 0, \dots, num_parts - 1\}$ where num_parts is the number of parts in our system. -1 denotes background and $0, \dots, num_parts - 1$ denote the indices for the part that is chosen.

It is important to note that while the edge map $E(t, f, e)$ is a coding over particular time-frequency regions whether a particular edge type is present, the feature map $\Phi(t, f)$ is a coding over regions of the edge map, in particular its over the region $[t, t + part_width) \times [f, f + part_height)$ and over all edge types. Since we model the edge map features as conditionally independent

bernoullis the coding follows the formula:

$$k \arg \max_{s \in [0, part_width], g \in [0, part_height], e} E(t+s, f+g, e) \log p_{k,s,g,e} + (1-E(t, f, e)) \log(1-p_{k,s,g,e}).$$

We only perform this feature transformation on patches with sufficient edge activity:

$$\sum_{s \in [0, part_width], g \in [0, part_height], e} E(t+s, f+g, e) \geq \tau_9,$$

i.e. time-frequency regions where the edge activity is greater than the 90th percentile. We can then think of a coding induced on the spectrogram by doing this. Namely, for non-background patch b with part code k whose root location is at (t, f) we construct a kernel η_b such that $\eta_b(t+s, f+g) = spec_patch_k$ where $spec_patch_k$ is the average over the implicitly clustered spectrogram clustering induced by the bernoulli model clustering.

```
import matplotlib.pyplot as plt
import matplotlib.cm as cm
import numpy as np
plt.figure()
plt.subplot(4,1,1)
S = np.load('../data/1S.npy')
plt.imshow(S[::-1], interpolation='nearest', aspect=3)
plt.xticks(np.arange(7)*100, tuple(str(i*100/5) for i in xrange(7)))
# since 8*16000 (sample rate)/512 (fft length) = 250
plt.yticks(np.arange(6)*8, tuple(str(i*250) for i in xrange(6))[:, -1])
plt.xlabel('time (ms)')
plt.ylabel('freq (Hz)')

for code_id, num_parts in enumerate([15, 20, 25]):
    plt.subplot(4,1, code_id+2)
    S_coded = np.load('../data/1S_coded%d.npy' % num_parts)
    plt.imshow(S_coded[::-1], cmap=cm.bone, interpolation='nearest', aspect=3)
    plt.title('S coded with %d parts' % num_parts)
    plt.xticks(np.arange(7)*100, tuple(str(i*100/5) for i in xrange(7)))
    # since 8*16000 (sample rate)/512 (fft length) = 250
    plt.yticks(np.arange(6)*8, tuple(str(i*250) for i in xrange(6))[:, -1])
    plt.xlabel('time (ms)')
```

```

plt.ylabel('freq (Hz)')

plt.subplots_adjust(hspace=1,left=0)
plt.savefig('1S_coded_compare.png')

plt.figure()
plt.subplot(4,1,1)
S = np.load('../data/1S.npy')
plt.imshow(S[::-1],interpolation='nearest',aspect=3)
plt.xticks(np.arange(7)*100,tuple(str(i*100/5) for i in xrange(7)))
# since 8*16000 (sample rate)/512 (fft length) = 250
plt.yticks(np.arange(6)*8,tuple(str(i*250) for i in xrange(6))[:,::-1])
plt.xlabel('time (ms)')
plt.ylabel('freq (Hz)')

for code_id,num_parts in enumerate([15,20,25]):
    plt.subplot(4,1,code_id+2)
    S_coded = np.load('../data/1S_coded%d.npy' % num_parts)
    plt.imshow(S_coded[::-1],interpolation='nearest',aspect=3)
    plt.title('S coded with %d parts' % num_parts)
    plt.xticks(np.arange(7)*100,tuple(str(i*100/5) for i in xrange(7)))
    # since 8*16000 (sample rate)/512 (fft length) = 250
    plt.yticks(np.arange(6)*8,tuple(str(i*250) for i in xrange(6))[:,::-1])
    plt.xlabel('time (ms)')
    plt.ylabel('freq (Hz)')

plt.subplots_adjust(hspace=1,left=0)
plt.savefig('1S_coded_compare_color.png')

```

The next step is to visualize what happens when we denoise or code a spectrogram using the parts. We take the estimated templates from the EM algorithm and we are going to use these to code spectrograms. The step is going




Figure 18: Spectrogram Part Clusters - 50 components




Figure 19: Spectrogram Part Clusters - 80 components




Figure 20: Spectrogram Part Clusters - 100 components




Figure 21: Spectrogram Part Clusters - 20 components, 1 Utterance

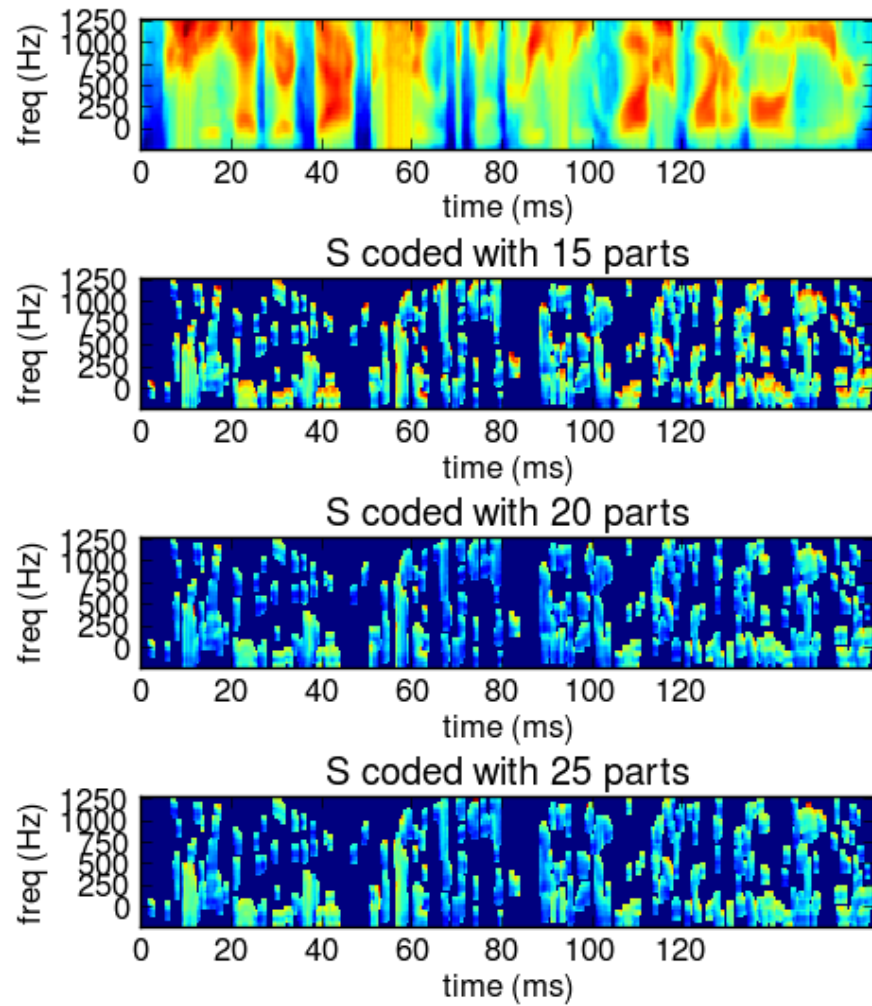


Figure 22: Spectrogram Coding Comparison